

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```

import nltk
import string

from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

import pickle
from tqdm import tqdm
import os
from collections import Counter

# ===== loading libraries =====
=====

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score

from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn import metrics

```

```

from sklearn.metrics import roc_curve, auc

from sklearn.feature_extraction.text import CountVectorizer

from prettytable import PrettyTable

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

import itertools

```

```

C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkiz
e to chunkize_serial
  warnings.warn("detected Windows; aliasing chunkize to chunkize_seria
l")

```

In [2]: `import os`
`print(os.listdir("."))`

```

['.ipynb_checkpoints', '04 Amazon Fine Food Reviews Analysis_NaiveBaye
s.ipynb', '05 Amazon Fine Food Reviews Analysis_Logistic Regression.ipy
nb', '05 Amazon Fine Food Reviews Analysis_Logistic Regression_origina
l.ipynb', 'Amazon Logistic Regression Submission.ipynb', 'Amazon_Logist
ic_Regression_Submission.ipynb', 'database.sqlite', 'Logistic Regressio
n in progress.ipynb', 'LogisticRegression.ipynb', 'LogisticRegression_c
opy.ipynb', 'Logistic_Regression_from_KNN_fun.ipynb', 'model_avgw2v_l1.
pkl', 'model_avgw2v_l2.pkl', 'preprocessed_final', 'sent_vectors_test.p
kl', 'sent_vectors_train.pkl', 'Some subtle python operations.ipynb']

```

In [3]: `# from google.colab import files`
`# files.upload()`

`# !pip install -q kaggle`

`# !mkdir -p ~/.kaggle`
`# !cp kaggle.json ~/.kaggle/`

`# !kaggle datasets download -d snap/amazon-fine-food-reviews`

```
# !unzip amazon-fine-food-reviews.zip
```

Data Import and Preprocessing

```
In [0]: # Google Drive
#final = pd.read_pickle('/content/drive/My Drive/Aaic/final_0317.pkl')
```

```
In [0]: # Local
#final = final = pickle.load(open('preprocessed_final', 'rb'))
```

```
In [8]: #!ls
```

```
amazon-fine-food-reviews.zip  hashes.txt  Reviews.csv
database.sqlite               kaggle.json sample_data
```

```
In [4]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 50
0000 data points
# you can change the number to any other number based on your computing
power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Sco
re != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points
```

```
In [5]: filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score
!= 3 LIMIT 100000""", con)
```

```
In [6]: # Give reviews with Score>3 a positive rating(1), and reviews with a sc
ore<3 a negative rating(0).
```

```
def partition(x):
    if x < 3:
        return 0
    return 1
```

```
In [7]: #changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (100000, 10)

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	

```
In [8]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
```

```
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [9]: print(display.shape)
display.head()
```

```
(80668, 7)
```

```
Out[9]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [10]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out[10]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
--	--------	-----------	-------------	------	-------	------	----------

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

```
In [11]: display['COUNT(*)'].sum()
```

```
Out[11]: 393063
```

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [12]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out[12]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenon
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenon
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [13]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

```
In [14]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

```
Out[14]: (87775, 10)
```

```
In [15]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[15]: 87.775
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [16]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
Out[16]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenom
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	

1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	
---	-------	------------	----------------	-----	---	--



In [17]: `final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]`

In [18]: `final.head()`

Out[18]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
22620	24750	2734888454	A13ISQV0U9GZIC	Sandikaye	1	
22621	24751	2734888454	A1C298ITT645B6	Hugh G. Pritchard	0	
70677	76870	B00002N8SM	A19Q006CSFT011	Arielle	0	

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessD
70676	76869	B00002N8SM	A1FYH4S02BW7FN	wonderer	0	
70675	76868	B00002N8SM	AUE8TB5VHS6ZV	eyeofthestorm	0	

```
In [19]: #Before starting the next phase of preprocessing lets see the number of
          entries left
          print(final.shape)

          #How many positive and negative reviews are present in our dataset?
          final['Score'].value_counts()

          (87773, 10)
```

```
Out[19]: 1    73592
          0    14181
          Name: Score, dtype: int64
```

Observation: This is an imbalance dataset. There are roughly 84% Positive review and 16% Negative reviews.

Checkpoint 1: Imported data and some basic eda

```
In [0]: # Filtering 1 (positive) and 0 (negative) reviews
          #pos_reviews = final[final['Score'] == 1].sample(n=10000, random_state=
          507)
          #neg_reviews = final[final['Score'] == 0].sample(n=10000, random_state=
          507)
```

```
# Combining the above dataframes into one
#final = pd.concat([pos_reviews, neg_reviews])
```

```
In [20]: # Converting Time to time format in seconds using a unix epoch time
# We will arrange the entire final dataframe in ascending order for time based splitting

final["Time"] = pd.to_datetime(final["Time"], origin='unix', unit = "s")
final = final.sort_values(by = "Time")
```

```
In [21]: final.size
```

```
Out[21]: 877730
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [22]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

=====

I have made these brownies for family and for a den of cub scouts and no one would have known they were gluten free and everyone asked for seconds! These brownies have a fudgy texture and have bits of chocolate chips in them which are delicious. I would say the mix is very thick and a little difficult to work with. The cooked brownies are slightly difficult to cut into very neat edges as the edges tend to crumble a little and I would also say that they make a slightly thinner layer of brownies than most of the store brand gluten containing but they taste just as good, if not better. Highly recommended!

(For those wondering, this mix requires 2 eggs OR 4 egg whites and 7 tbs melted butter to prepare. They do have suggestions for lactose free and low fat preparations)

=====

This gum is my absolute favorite. By purchasing on amazon I can get the savings of large quantities at a very good price. I highly recommend to all gum chewers. Plus as you enjoy the peppermint flavor and freshening of breath you are whitening your teeth all at the same time.

=====

This is an excellent product, both tasty and priced right. It's difficult to find this product in regular local grocery stores, so I was thrilled to find it.

=====

```
In [23]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

```
In [24]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)
```

```
soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

=====

I have made these brownies for family and for a den of cub scouts and no one would have known they were gluten free and everyone asked for seconds! These brownies have a fudgy texture and have bits of chocolate chips in them which are delicious. I would say the mix is very thick and a little difficult to work with. The cooked brownies are slightly difficult to cut into very neat edges as the edges tend to crumble a little and I would also say that they make a slightly thinner layer of brownies than most of the store brand gluten containing but they taste just as good, if not better. Highly recommended!(For those wondering, this mix requires 2 eggs OR 4 egg whites and 7 tbs melted butter to prepare. They do have suggestions for lactose free and low fat preparations)

=====

This gum is my absolute favorite. By purchasing on amazon I can get the savings of large quantities at a very good price. I highly recommend to all gum chewers. Plus as you enjoy the peppermint flavor and freshening of breath you are whitening your teeth all at the same time.

=====

This is an excellent product, both tasty and priced right. It's difficult to find this product in regular local grocery stores, so I was thrilled to find it.

In [25]: `# https://stackoverflow.com/a/47091490/4084039`
`import re`


```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [26]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

This gum is my absolute favorite. By purchasing on amazon I can get the savings of large quantities at a very good price. I highly recommend to all gum chewers. Plus as you enjoy the peppermint flavor and freshening of breath you are whitening your teeth all at the same time.

=====

```
In [27]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

```
In [28]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

This gum is my absolute favorite By purchasing on amazon I can get the savings of large quantities at a very good price I highly recommend to a ll gum chewers Plus as you enjoy the peppermint flavor and freshing of breath you are whitening your teeth all at the same time

```
In [29]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'no
t'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in
the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'o
urs', 'ourselves', 'you', "you're", "you've", \
               "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselve
s', 'he', 'him', 'his', 'himself', \
               'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'it
s', 'itself', 'they', 'them', 'their', \
               'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'th
is', 'that', "that'll", 'these', 'those', \
               'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'h
ave', 'has', 'had', 'having', 'do', 'does', \
               'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',
'because', 'as', 'until', 'while', 'of', \
               'at', 'by', 'for', 'with', 'about', 'against', 'between',
'into', 'through', 'during', 'before', 'after', \
               'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out',
'on', 'off', 'over', 'under', 'again', 'further', \
               'then', 'once', 'here', 'there', 'when', 'where', 'why', 'h
ow', 'all', 'any', 'both', 'each', 'few', 'more', \
               'most', 'other', 'some', 'such', 'only', 'own', 'same', 's
o', 'than', 'too', 'very', \
               's', 't', 'can', 'will', 'just', 'don', "don't", 'should',
"should've", 'now', 'd', 'll', 'm', 'o', 're', \
```

```
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't",
'didn', "didn't", 'doesn', "doesn't", 'hadn', \
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "is
n't", 'ma', 'mightn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
"shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

```
In [30]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower
() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 87773/87773
[00:30<00:00, 2909.84it/s]
```

```
In [31]: preprocessed_reviews[1500]
```

```
Out[31]: 'gum absolute favorite purchasing amazon get savings large quantities go
od price highly recommend gum chewers plus enjoy peppermint flavor fres
hing breath whitening teeth time'
```

[3.2] Preprocessing Review Summary

```
In [32]: ## Similarly you can do preprocessing for review summary also.
from tqdm import tqdm
```

```

preprocessed_summary = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Summary'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower
    () not in stopwords)
    preprocessed_summary.append(sentence.strip())

```

```

2%|██████████| 1694/87773
[00:00<00:32, 2682.41it/s]C:\Users\Nit-prj1010\AppData\Local\Continuum
\anaconda3\lib\site-packages\bs4\__init__.py:273: UserWarning: "b'...'"
looks like a filename, not markup. You should probably open this file a
nd pass the filehandle into BeautifulSoup.
' BeautifulSoup.' % markup)
43%|██████████| 37757/87773
[00:09<00:11, 4525.80it/s]C:\Users\Nit-prj1010\AppData\Local\Continuum
\anaconda3\lib\site-packages\bs4\__init__.py:273: UserWarning: "b'...'"
looks like a filename, not markup. You should probably open this file a
nd pass the filehandle into BeautifulSoup.
' BeautifulSoup.' % markup)
56%|██████████| 48811/87773
[00:12<00:11, 3520.87it/s]C:\Users\Nit-prj1010\AppData\Local\Continuum
\anaconda3\lib\site-packages\bs4\__init__.py:273: UserWarning: "b'...'"
looks like a filename, not markup. You should probably open this file a
nd pass the filehandle into BeautifulSoup.
' BeautifulSoup.' % markup)
82%|██████████| 72288/87773
[00:17<00:02, 5350.16it/s]C:\Users\Nit-prj1010\AppData\Local\Continuum
\anaconda3\lib\site-packages\bs4\__init__.py:273: UserWarning: "b'...'"
looks like a filename, not markup. You should probably open this file a
nd pass the filehandle into BeautifulSoup.

```

[illegible]

```
In [33]: final.head()
```

Out[33]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessD
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar		0
1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie		7
1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase		10
28086	30629	B00008RCMI	A19E94CF5O1LY7	Andrew Arnold		0
28087	30630	B00008RCMI	A284C7M23F0APC	A. Mendoza		0

```
In [34]: final['CleanedText']= preprocessed_reviews  
final['CleanedSummary']= preprocessed_summary
```

```
In [35]: final.CleanedText.isnull().sum()
```

```
Out[35]: 0
```

```
In [36]: final.head()
```

```
Out[36]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessD
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	0	
1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	
1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	
28086	30629	B00008RCMI	A19E94CF5O1LY7	Andrew Arnold	0	

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessD
	28087	30630	B00008RCMI	A284C7M23F0APC	A. Mendoza	0

```
In [0]: # writing the preprocessed final dataframe to dist
#file_Name = "preprocessed_final"
# open the file for writing
#fileObject = open(file_Name,'wb')

# this writes the object a to the
# file named 'testfile'
#pickle.dump(final,fileObject)
```

Load saved final

```
In [0]: #final = pickle.load(open('preprocessed_final', 'rb'))
```

```
In [37]: final.columns
```

```
Out[37]: Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
               'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text',
               'CleanedText', 'CleanedSummary'],
              dtype='object')
```

Checkpoint 2: Data is now sorted based on Time and preprocessed.

```
In [38]: # Create X and Y variable
X = final['CleanedText'].values
y= final['Score'].values
```

```
In [39]: type(X)
```

```
Out[39]: numpy.ndarray
```

```
In [40]: type(y)
```

```
Out[40]: numpy.ndarray
```

```
In [41]: # ss
         from sklearn.model_selection import train_test_split

         # Splitting into train and test in the ratio 70:30
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3
0,shuffle=False, random_state=507)
         #X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test
_size=0.30, shuffle=False, random_state=507)
```

```
In [42]: print("Train Set:",X_train.shape, y_train.shape[0])
         print("Test Set:",X_test.shape, y_test.shape[0])
```

```
Train Set: (61441,) 61441
```

```
Test Set: (26332,) 26332
```

Checkpoint 3: Data has been partitioned into train, cv and test

Defining functions that we will be using throughout the notebook for BoW, TFIDF, AvgW2V, TFIDF-WW2V

1. **get_best_hyperparameter_C(vectorizer, X_train, X_test, y_train, y_test, penalty)** : This function will run GridSearchCV with cv = 5 and the penalty on the training and test set data specified by the user.
2. **plot_auc(model, X_train, X_test)**: This function will plot the AUC curve
3. **most_informative_feature_for_binary_classification(vectorizer, classifier, n=10)** : This function will return the most important features for the positive and the negative class
4. **print_confusion_matrix(model, X_train, X_test)**: Prints the confusion matrix for the train and test set data.

5. `plot_confusion_matrix_heatmap(cm, classes, normalize=False, title='Confusion matrix', cmap=plt.cm.Blues)` : Taken from the official sklearn website, this function will return heatmap representation of the confusion matrix
6. `plot_heatmap_confusion_matrix(model, X_test)`: Calculates the confusion matrix on the dataset provided and passes it to `plot_confusion_matrix_heatmap()` to print the heatmap

Finding the hyper parameter C (1/lambda) using RandomSearchCV with cv = 5

```
In [43]: def get_best_hyperparameter_C(vectorizer, X_train, X_test, y_train, y_test, penalty_l):

    """
    This function takes in the vectorizer, and performs LogisticRegression hyperparameter tuning using GridSearchCV with 5 fold cv
    Returns the value of hyperparameter C and draws the error plot for various values of C

    Usage: get_best_hyperparameter_C(vectorizer, X_train, X_test, y_train, y_test, penalty)
    """
    tuned_parameters = [{ 'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}]
    alpha = [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4] #k

    #tuned_parameters = [{ 'C': [10**-4, 10**-3, 10**-2]}]
    #alpha = [10**-4, 10**-3, 10**-2]

    # Using GridSearchCVSearchCV with 5 fold cv
    gs_obj = GridSearchCV(LogisticRegression(penalty= penalty_l), tuned_parameters, scoring = 'roc_auc', cv=5)

    gs_obj.fit(X_train, y_train)

    train_auc= gs_obj.cv_results_['mean_train_score']
    train_auc_std= gs_obj.cv_results_['std_train_score']
```

```

cv_auc = gs_obj.cv_results_['mean_test_score']
cv_auc_std= gs_obj.cv_results_['std_test_score']

# draws the error plot

plt.plot(alpha, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

plt.plot(alpha, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
plt.legend()
plt.xlabel("log(C)- hyperparamter")
plt.xscale('log')
plt.ylabel("AUC")
plt.title("ERROR PLOT")
plt.show()

# Results of the gs object

# Code https://stackoverflow.com/questions/42793254/what-replaces-gridsearchcv-grid-scores-in-scikit#answer-42800056
means = gs_obj.cv_results_['mean_test_score']
stds = gs_obj.cv_results_['std_test_score']

t1 = PrettyTable()
t1.field_names = ['Mean CV Score', 'Std CV Score', 'Param']

for mean, std, params in zip(means, stds, gs_obj.cv_results_['params']):

```

```

        t1.add_row([round(mean, 3), round(std * 2,5), params])

    print(t1)

    print("\nThe best estimator:{}".format(gs_obj.best_estimator_))
    print("\nThe best score is:{}".format(gs_obj.best_score_))
    print("The best value of C is:{}".format(gs_obj.best_params_))

    # Returns the mean accuracy on the given test data and labels.
    print("Mean Score: {}".format(gs_obj.score(X_test, y_test)))

    return gs_obj.best_params_

```

train and test AUC

```

In [44]: def plot_auc(model, X_train, X_test):

        """
        This function will plot the AUC for the vectorized train and test d
        ata.
        Returns the plot and also the values of auc for train and test

        Usage: auc_train, auc_test = plot_auc(model, X_train, X_test)
        """

        train_fpr, train_tpr, thresholds = roc_curve(y_train, model.predict
        _proba(X_train)[:,-1])
        test_fpr, test_tpr, thresholds = roc_curve(y_test, model.predict_pr
        oba(X_test)[:,-1])

        plt.plot([0,1],[0,1], 'k--')
        plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fp
        r, train_tpr)))
        plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, t
        est_tpr)))
        plt.legend()
        plt.xlabel("fpr")

```

```
plt.ylabel("tpr")
plt.title("ROC Curve")
plt.show()

print("train AUC: {}".format(auc(train_fpr, train_tpr)))
print("test AUC: {}".format(auc(test_fpr, test_tpr)))

return auc(train_fpr, train_tpr), auc(test_fpr, test_tpr)
```

important features

```
In [45]: # https://stackoverflow.com/questions/26976362/how-to-get-most-informative-features-for-scikit-learn-classifier-for-different-c
def most_informative_feature_for_binary_classification(vectorizer, classifier, n=10):

    """
        Takes in the vectorizer, classifier (model) and the number of important features to return

        Usage: most_informative_feature_for_binary_classification(vectorizer, classifier, n=10)
    """
    class_labels = classifier.classes_
    feature_names = vectorizer.get_feature_names()
    topn_class_0 = sorted(zip(classifier.coef_[0], feature_names))[:n]
    topn_class_1 = sorted(zip(classifier.coef_[0], feature_names))[-n:]

    t1 = PrettyTable()
    t1.field_names = ['Class', 'Coefficient (Importance)', 'Feature Name']

    for coef, feat in topn_class_0:
        t1.add_row([class_labels[0], abs(coef), feat])

    print(t1)
```

```

print("""*52)

t2 = PrettyTable()
t2.field_names = ['Class', 'Coefficient (Importance)', 'Feature Name']

for coef, feat in reversed(topn_class_1):
    t2.add_row([class_labels[1], abs(coef), feat])

print(t2)

#for coef, feat in topn_class1:
#    if coef < 0:
#        print(class_labels[0], abs(coef), feat)

#print("""*30)

#for coef, feat in reversed(topn_class2):
#    if coef > 0:
#        print(class_labels[1], abs(coef), feat)

```

print confusion matrix

```

In [46]: def print_confusion_matrix(model, X_train, X_test):
        """
        Takes in the model, X_train, X_test and prints the confusion matrix
        Usage: print_confusion_matrix(model, X_train, X_test)
        """
        print("*****Train confusion matrix*****")
        print(confusion_matrix(y_train, model.predict(X_train)))
        print("\n*****Test confusion matrix*****")
        print(confusion_matrix(y_test, model.predict(X_test)))

```

heat map of confusion matrix

```

In [47]: # Code modified from sklearn tutorial: https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_confusion\_matrix.html

# Heat map of confusion matrix

def plot_confusion_matrix_heatmap(cm, classes,
                                  normalize=False,
                                  title='Confusion matrix',
                                  cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    #if normalize:
    #    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    #    print("Normalized confusion matrix")
    #else:
    #    print('Confusion matrix')

    #print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()

```

```
In [48]: # def plot_heatmap_confusion_matrix(model, X_test):
#         cnf_matrix = confusion_matrix(y_test, model.predict(X_test))
#         np.set_printoptions(precision=2)
#         class_names = ['Negative', 'Positive']
#         # Plot non-normalized confusion matrix
#         plt.figure()
#         plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, ti
tle='Test Set Confusion Matrix');
```

```
In [49]: # def plot_heatmap_confusion_matrix_train(model, X_train):
#         cnf_matrix = confusion_matrix(y_train, model.predict(X_train))
#         np.set_printoptions(precision=2)
#         class_names = ['Negative', 'Positive']
#         # Plot non-normalized confusion matrix
#         plt.figure()
#         plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, ti
tle='Train Set Confusion Matrix');
```

[4.1] BAG OF WORDS

```
In [50]: # ss
from sklearn.feature_extraction.text import CountVectorizer
bow_vectorizer= CountVectorizer(ngram_range=(1,2), min_df=10, max_featu
res=10000)
bow_vectorizer.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_bow = bow_vectorizer.transform(X_train)
#X_cv_bow = vectorizer.transform(X_cv)
X_test_bow = bow_vectorizer.transform(X_test)

print("After vectorizations")
print(X_train_bow.shape, y_train.shape)
#print(X_cv_bow.shape, y_cv.shape)
```

```
print(X_test_bow.shape, y_test.shape)
print("="*100)
```

```
After vectorizations
(61441, 10000) (61441,)
(26332, 10000) (26332,)
```

```
=====
=====
```

```
In [51]: print("the type of count vectorizer ",type(X_train_bow))
print("the shape of cut text BOW vectorizer ",X_train_bow.get_shape())
print("the number of unique words: ", X_train_bow.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of cut text BOW vectorizer (61441, 10000)
the number of unique words: 10000
```

Standardize the data

```
In [52]: # We will set the attribute with_mean = False, as StandardScaler does not work on sparse matrix
# when attempted on sparse matrices, because centering them entails building a dense matrix which in common use cases
# is likely to be too large to fit in memory. ---> sklearn documentation
```

```
from sklearn.preprocessing import StandardScaler
X_train_bow=StandardScaler(with_mean=False).fit_transform(X_train_bow)
X_test_bow=StandardScaler(with_mean=False).fit_transform(X_test_bow)

print(X_train_bow.shape, y_train.shape)

print(X_test_bow.shape, y_test.shape)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
```

```
warnings.warn(msg, DataConversionWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
```



```

C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\utils\validation.py:595: DataConversionWarning: Data with inp
ut dtype int64 was converted to float64 by StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\utils\validation.py:595: DataConversionWarning: Data with inp
ut dtype int64 was converted to float64 by StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\utils\validation.py:595: DataConversionWarning: Data with inp
ut dtype int64 was converted to float64 by StandardScaler.
  warnings.warn(msg, DataConversionWarning)

```

```

(61441, 10000) (61441,)
(26332, 10000) (26332,)

```

[5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

```

In [107]: # Free up some space
del final, display, preprocessed_reviews, preprocessed_summary, X, y, so
rted_data, filtered_data, actualScore, sent_0, sent_1000, sent_150, sen
t_1500, sent_4900
#del sent_vectors_cv, sent_vectors_test, sent_vectors_train
#del tfidf_sent_vectors_train, tfidf_sent_vectors_cv, tfidf_sent_vector
s_test
#del list_of_sentence_cv, list_of_sentence_test, list_of_sentence_train
#del dictionary, stopwords, tfidf_feat, w2v_words

```

```

In [53]: # Vectorizer = BoW, penalty = l1
best_estimator_bow_l1= get_best_hyperparameter_C(bow_vectorizer, X_train
_bow, X_test_bow, y_train, y_test, penalty_l = 'l1' )

```

```

C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
  "the number of iterations.", ConvergenceWarning)


```

```

C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package

```

```
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)
```



Mean CV Score	Std CV Score	Param
0.5	0.0	{'C': 0.0001}
0.857	0.02148	{'C': 0.001}
0.946	0.00967	{'C': 0.01}
0.931	0.01305	{'C': 0.1}
0.906	0.01855	{'C': 1}
0.899	0.01977	{'C': 10}
0.89	0.02458	{'C': 100}
0.883	0.02402	{'C': 1000}
0.88	0.02441	{'C': 10000}

```
The best estimator:LogisticRegression(C=0.01, class_weight=None, dual=False,
fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l1', random_state=None, solver='warn',
```

```
tol=0.0001, verbose=0, warm_start=False)
```


```
The best score is:0.945815885045128  
The best value of C is:{'C': 0.01}  
Mean Score: 0.9547743986302348
```

```
In [66]: # https://stackoverflow.com/questions/3097866/access-an-arbitrary-element-in-a-dictionary-in-python  
#list(my_dict.keys())[0]  
list(best_estimator_bow_l1.values())[0]
```

```
Out[66]: 0.01
```

```
In [54]: model_bow_l1 = LogisticRegression(C= list(best_estimator_bow_l1.values  
())[0] ,penalty = 'l1')  
model_bow_l1.fit(X_train_bow,y_train)  
y_pred = model_bow_l1.predict(X_test_bow)
```

```
In [55]: # AUC-ROC plot  
auc_train_bow_l1, auc_test_bow_l1 = plot_auc(model_bow_l1, X_train_bow,  
X_test_bow)
```



```
train AUC: 0.9695167883807346
test AUC: 0.9547676364720178
```

```
In [56]: # Confusion Matrix
print_confusion_matrix(model_bow_l1, X_train_bow, X_test_bow)
```

```
*****Train confusion matrix*****
[[ 6280  3344]
 [   655 51162]]


*****Test confusion matrix*****
[[ 2785  1772]
 [   376 21399]]
```

```
In [57]: plt.figure(1)
plt.figure(figsize=(15, 4))

plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_bow_l1.predict(X_test_bow))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
est Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_bow_l1.predict(X_train_bow
))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
rain Set Confusion Matrix');
```

```
<Figure size 432x288 with 0 Axes>
```



Observation

1. For the BoW vectorizer, we calculated $C = 0.01$ using GridSearchCV with $cv = 5$ and with penalty l1.
2. We got train AUC: 0.9695167883807346 and test AUC: 0.9547676364720178
3. Using the confusion matrix, we can say that our model correctly predicted 21399 positive reviews and 2785 negative reviews.
4. The model incorrectly classified 376 negative reviews and 1772 positive reviews.
5. The True Postive Rate is 98.27 and the True Negative Rate is 61.11
6. The accuracy of the model is 91.84

```
In [58]: # Confustion Matrix heatmap
# print("Train set")
# plot_heatmap_confusion_matrix(model_bow_l1, X_train_bow)
# plot_heatmap_confusion_matrix_train(model_bow_l1, X_train_bow)
# confusion_matrix(y_train, model_bow_l1.predict(X_train_bow))

# Confustion Matrix heatmap
# print("Train set")
# plot_heatmap_confusion_matrix(model_bow_l1, X_test_bow)
```

[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW,

SET 1

```
In [59]: # The model has already been fitted. Here we are just going to calculate the sparsity on weight vector
w = model_bow_l1.coef_
print("The sparsity (no. of non-zero elements) in weight vector is {}".format(np.count_nonzero(w)))
```

The sparsity (no. of non-zero elements) in weight vector is 1910

Feature Engineering Let us perform FE to see if we can further improve the model. Here, we will append length of reviews as another feature.

```
In [60]: def get_text_length(x):
        """
        This function takes in a array and returns the length of the elements in the array.
        """
        return np.array([len(t) for t in x]).reshape(-1, 1)
```

```
In [61]: rev_len_X_train = get_text_length(X_train)
rev_len_X_test = get_text_length(X_test)
```

```
In [62]: from sklearn.feature_extraction.text import CountVectorizer
bow_vectorizer_fe = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=10000)
bow_vectorizer_fe.fit(X_train) # fit has to happen only on train data
```

```
Out[62]: CountVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=10000, min_df=10,
ngram_range=(1, 2), preprocessor=None, stop_words=None,
strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, vocabulary=None)
```

```
In [63]: # we use the fitted CountVectorizer to convert the text to vector
X_train_bow = bow_vectorizer_fe.transform(X_train)
```

```
X_test_bow = bow_vectorizer_fe.transform(X_test)
```

```
print("After vectorizations")
print(X_train_bow.shape, y_train.shape)
print(X_test_bow.shape, y_test.shape)
print("="*100)
```

```
After vectorizations
(61441, 10000) (61441,)
(26332, 10000) (26332,)
```

```
=====
=====
```

Standardize the data

```
In [64]: # We will set the attribute with_mean = False, as StandardScaler does not
work on sparse matrix
# when attempted on sparse matrices, because centering them entails building a dense matrix which in common use cases
# is likely to be too large to fit in memory. ---> sklearn documentation

from sklearn.preprocessing import StandardScaler
X_train_bow=StandardScaler(with_mean=False).fit_transform(X_train_bow)
X_test_bow=StandardScaler(with_mean=False).fit_transform(X_test_bow)

print(X_train_bow.shape, y_train.shape)

print(X_test_bow.shape, y_test.shape)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\utils\validation.py:595: DataConversionWarning: Data with inp
ut dtype int64 was converted to float64 by StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\utils\validation.py:595: DataConversionWarning: Data with inp
ut dtype int64 was converted to float64 by StandardScaler.
  warnings.warn(msg, DataConversionWarning)

(61441, 10000) (61441,)
(26332, 10000) (26332,)
```

```
In [65]: type(rev_len_X_train)
```


```
Out[65]: numpy.ndarray
```

```
In [66]: type(X_train_bow)
```

```
Out[66]: scipy.sparse.csr.csr_matrix
```

```
In [67]: from scipy.sparse import hstack
# Here we append the sparse matrix and the dense array that contains th
e length of the text passed to it
X_train_bow_fe = hstack((X_train_bow, np.array(rev_len_X_train)))
X_test_bow_fe = hstack((X_test_bow, np.array(rev_len_X_test)))
```

```
In [68]: # Get the best hyperparameter using GridSearchCV with penalty l1 and cv
= 5
best_estimator_bow_l1_fe = get_best_hyperparameter_C(bow_vectorizer_fe,
X_train_bow_fe, X_test_bow_fe, y_train, y_test, penalty_l = 'l1' )
```

Mean CV Score	Std CV Score	Param
0.44	0.01866	{'C': 0.0001}
0.849	0.02492	{'C': 0.001}
0.945	0.00919	{'C': 0.01}
0.931	0.01316	{'C': 0.1}
0.906	0.01843	{'C': 1}
0.898	0.0208	{'C': 10}
0.888	0.02375	{'C': 100}
0.882	0.02482	{'C': 1000}
0.881	0.02594	{'C': 10000}

The best estimator:LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l1', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)


The best score is:0.9450663381665366

The best value of C is:{'C': 0.01}

Mean Score: 0.9538766994570484

```
In [69]: # Fitting the BoW vectorizer on LogisticRegression Model with penalty l
1 and C = 0.01
model_bow_l1_fe = LogisticRegression(C= list(best_estimator_bow_l1_fe.v
alues())[0] ,penalty = 'l1')
model_bow_l1_fe.fit(X_train_bow_fe,y_train)
y_pred = model_bow_l1_fe.predict(X_test_bow_fe)
```

```
In [70]: # AUC-ROC plot
auc_train_bow_l1_fe, auc_test_bow_l1_fe = plot_auc(model_bow_l1_fe, X_t
rain_bow_fe, X_test_bow_fe)
```



```
train AUC: 0.9694355169306984
test AUC: 0.9538764071978185
```


```
In [71]: # Confusion Matrix
print_confusion_matrix(model_bow_l1_fe, X_train_bow_fe, X_test_bow_fe)

*****Train confusion matrix*****
[[ 6274  3350]
 [   652 51165]]
```

```
*****Test confusion matrix*****  
[[ 2788 1769]  
 [  362 21413]]
```

```
In [72]: # Confusion Matrix heatmap  
plt.figure(1)  
plt.figure(figsize=(15, 4))  
  
plt.subplot(121) # Test confusion matrix  
cnf_matrix = confusion_matrix(y_test, model_bow_l1_fe.predict(X_test_bo  
w_fe))  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T  
est Set Confusion Matrix');  
  
plt.subplot(122) # Train Confusion matrix  
cnf_matrix = confusion_matrix(y_train, model_bow_l1_fe.predict(X_train_  
bow_fe))  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T  
rain Set Confusion Matrix');
```

<Figure size 432x288 with 0 Axes>



Observation

1. For the BoW vectorizer with Feature Engineering, we calculated $C = 0.01$ using GridSearchCV with $cv = 5$ and with penalty l1.
2. We got train AUC: 0.9694354808358998 and test AUC: 0.9538772738827763
3. Using the confusion matrix, we can say that our model correctly predicted 21413 positive reviews and 2788 negative reviews.
4. The model incorrectly classified 362 negative reviews and 1769 positive reviews.
5. The True Postive Rate is 98.33 and the True Negative Rate is 61.18
6. The overall accuracy of the model is 91.90
7. **Doing Feature Engineering has made the model slightly perform better than the model without feature engineering.**

Observation

[5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

```
In [73]: # Vectorizer = BoW, penalty = l2
best_estimator_bow_l2= get_best_hyperparameter_C(bow_vectorizer, X_train)
```

```
n_bow, X_test_bow, y_train, y_test, penalty_l = 'l2' )
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```


```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv  
erge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package  
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
```

```
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Nit-prj1010\AppData\Local\Continuum\anaconda3\lib\site-package
s\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to conv
erge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
```



Mean CV Score	Std CV Score	Param
0.929	0.02038	{'C': 0.0001}
0.937	0.01645	{'C': 0.001}
0.921	0.01764	{'C': 0.01}
0.905	0.02018	{'C': 0.1}
0.898	0.02079	{'C': 1}
0.903	0.01494	{'C': 10}
0.905	0.02151	{'C': 100}
0.902	0.03122	{'C': 1000}
0.896	0.02433	{'C': 10000}


The best estimator: LogisticRegression(C=0.001, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)

The best score is: 0.9367812057739767
The best value of C is: {'C': 0.001}

Mean Score: 0.9415796492294188

```
In [74]: # Fitting the model with the best hyperparameter
model_bow_l2 = LogisticRegression(C= list(best_estimator_bow_l2.values
())[0] ,penalty = 'l2')
model_bow_l2.fit(X_train_bow,y_train)
y_pred = model_bow_l2.predict(X_test_bow)
```

```
In [75]: # AUC- ROC plot
auc_train_bow_l2, auc_test_bow_l2 = plot_auc(model_bow_l2, X_train_bow,
X_test_bow)
```



```
train AUC: 0.99309054211035
test AUC: 0.9415796492294188
```


```
In [76]: # Confusion Matrix
print_confusion_matrix(model_bow_l2, X_train_bow, X_test_bow)

*****Train confusion matrix*****
[[ 8166  1458]
 [  412 51405]]
```



```
*****Test confusion matrix*****  
[[ 3150  1407]  
 [   623 21152]]
```

```
In [77]: # Heatmap Confusion Matrix  
plt.figure(1)  
plt.figure(figsize=(15, 4))  
  
plt.subplot(121) # Test confusion matrix  
cnf_matrix = confusion_matrix(y_test, model_bow_l2.predict(X_test_bow))  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T  
est Set Confusion Matrix');  
  
plt.subplot(122) # Train Confusion matrix  
cnf_matrix = confusion_matrix(y_train, model_bow_l2.predict(X_train_bow  
)  
)  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T  
rain Set Confusion Matrix');  
  
<Figure size 432x288 with 0 Axes>
```



Observation

1. For the BoW vectorizer, we calculated $C = 0.001$ using GridSearchCV with $cv = 5$ and with penalty l2.
2. We got train AUC: 0.99309054211035 and test AUC: 0.9415796492294188
3. Using the confusion matrix, we can say that our model correctly predicted 21152 positive reviews and 3150 negative reviews.
4. The model incorrectly classified 623 negative reviews and 1407 positive reviews.
5. The True Postive Rate is 98.67 and the True Negative Rate is 60.03
6. The accuracy of the model is 91.98

```
In [0]: import sys
def sizeof_fmt(num, suffix='B'):
    ''' By Fred Cirera, after https://stackoverflow.com/a/1094933/1870254'''
    for unit in ['', 'Ki', 'Mi', 'Gi', 'Ti', 'Pi', 'Ei', 'Zi']:
        if abs(num) < 1024.0:
            return "%3.1f%s%s" % (num, unit, suffix)
        num /= 1024.0
    return "%.1f%s%s" % (num, 'Yi', suffix)

for name, size in sorted(((name, sys.getsizeof(value)) for name, value i
n locals()).items()),
```

```
key= lambda x: -x[1][:10]:  
print("{:>30}: {:>8}".format(name, sizeof_fmt(size)))
```

Perturbation Test

[5.1.2.1] Performing perturbation test (multicollinearity check) on BOW, SET 1

```
In [78]: # Step 1: Get the weights W after fitting your model with the data X  
W_vect = model_bow_l2.coef_
```

```
In [79]: # Step 2: Add a noise to the X ( $X' = X + e$ ) and get the new dataset  $X'$   
  
# We need to preserve our existing X and create a copy of it before we  
# can add noise to it  
# source: https://www.geeksforgeeks.org/copy-python-deep-copy-shallow-copy/  
# deepcopy(): any changes made to a copy of object do not reflect in the  
# original object.  
  
import copy  
X_train_bow_noisy = copy.deepcopy(X_train_bow) # copy of X_train_bow  
uniform_noise = np.random.uniform(-0.0001, 0.0001, 1) # adding some random  
uniform noise
```

```
In [80]: type(X_train_bow)
```

```
Out[80]: scipy.sparse.csr.csr_matrix
```

```
In [81]: # Since X_train_bow is sparse matrix, we do X.data += e  
X_train_bow_noisy.data += uniform_noise  
print(X_train_bow_noisy.shape)  
  
(61441, 10000)
```

```
In [82]: # Step 3: We fit the model again on data X_train_bow_noisy and get the  
weights W'
```

```
# Fitting the BoW vectorizer on LogisticRegression Model
model_bow_l2_noisy = LogisticRegression(C= list(best_estimator_bow_l2.v
alues())[0],penalty = 'l2')
model_bow_l2_noisy.fit(X_train_bow_noisy,y_train)
y_pred = model_bow_l2_noisy.predict(X_test_bow)
```

```
In [83]: # Get the weights W after fitting your model with the noisy data X'
W_vect_noisy = model_bow_l2_noisy.coef_
```

```
In [84]: # Step 4: Add small epsilon value to eliminate the 'divisible by zero'
error to W_vect and W_vect_noisy
esp = 10**-6
W_vect += esp
W_vect_noisy += esp
```

```
In [85]: # Step 5: Find % change between W_vect and W_vect_noisy

W_percent_change_vector = abs((W_vect - W_vect_noisy) / W_vect) *100
```

```
In [86]: W_percent_change_vector
```

```
Out[86]: array([[9.93e-04, 7.01e-06, 1.03e-02, ..., 2.59e-03, 5.13e-04, 6.51e-0
3]])
```

```
In [87]: # Step 6: Calculate 0th, 10th, 20th ..100th Percentile
percentile_range= range(0,101,10)
tab = PrettyTable()
tab.field_names = ['Percentile ', 'Percentile Value']

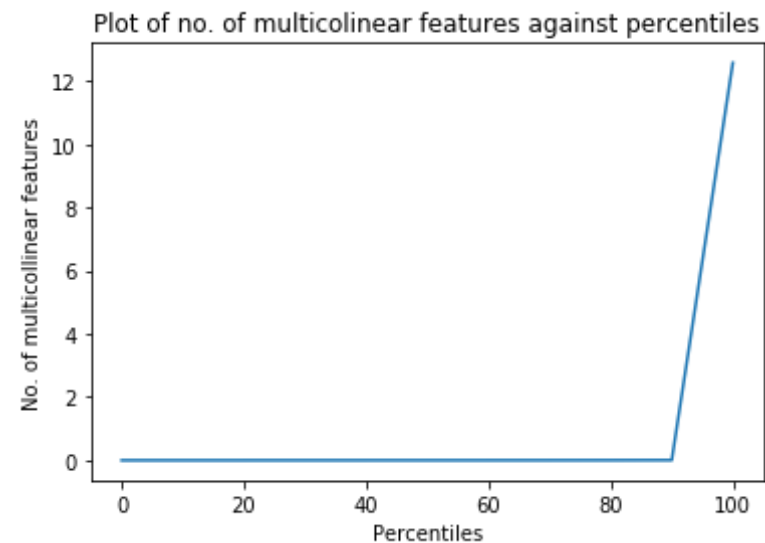
for p in percentile_range:
    tab.add_row([p, np.percentile(W_percent_change_vector,p)])

plt.plot(percentile_range, np.percentile(W_percent_change_vector,percen
tile_range))
```

```
plt.xlabel('Percentiles')
plt.ylabel('No. of multicollinear features')
plt.title("Plot of no. of multicollinear features against percentiles")

print(tab)
del (tab)
```

Percentile	Percentile Value
0	2.1288236710142e-08
10	0.0001290710220424033
20	0.00026651132954525614
30	0.000404709561975814
40	0.000581215567616487
50	0.0007999871796000233
60	0.0010965094631658796
70	0.0015858030927552692
80	0.002512170755860659
90	0.005024878518104683
100	12.572261483067537



Observation: We see that there is a sudden rise in values after the 90th percentile value.

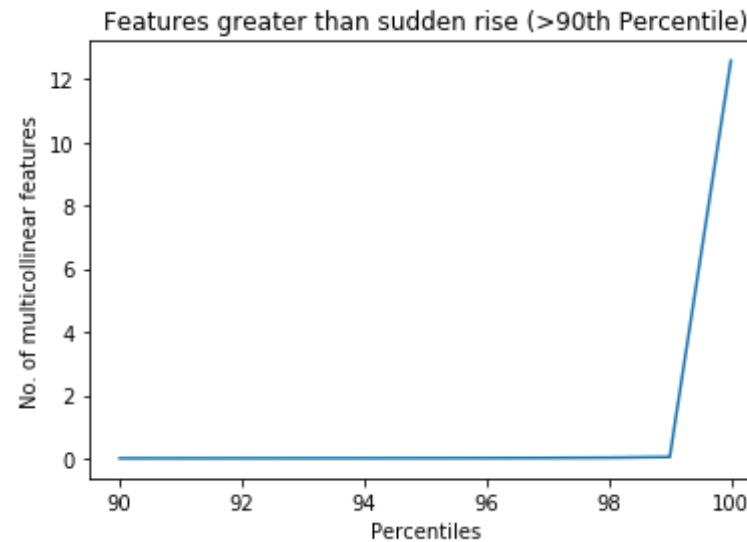
```
In [88]: # Plot values after the sudden rise i.e. from the 90th percentile value
s

sudden_rise= range(90,101,1)
tab = PrettyTable()
tab.field_names = ['Percentile ', 'Percentile Value']

for i in sudden_rise:
    tab.add_row([i, np.percentile(W_percent_change_vector,i)])

print(tab)
plt.plot(sudden_rise,np.percentile(W_percent_change_vector,sudden_rise)
);
plt.xlabel('Percentiles')
plt.ylabel('No. of multicollinear features')
plt.title("Features greater than sudden rise (>90th Percentile)")
del (tab)
```

Percentile	Percentile Value
90	0.005024878518104683
91	0.005536440548611815
92	0.006223372849219215
93	0.007020458291498725
94	0.008347083172674468
95	0.00971386294287677
96	0.011739505021271607
97	0.016008687942597848
98	0.024976125423512584
99	0.047399636184165865
100	12.572261483067537



Observation: Here, after 99th percentile, we see there is sudden rise in values.

In [89]: *# Plot values after the sudden rise i.e. from the 90th percentile value*
S

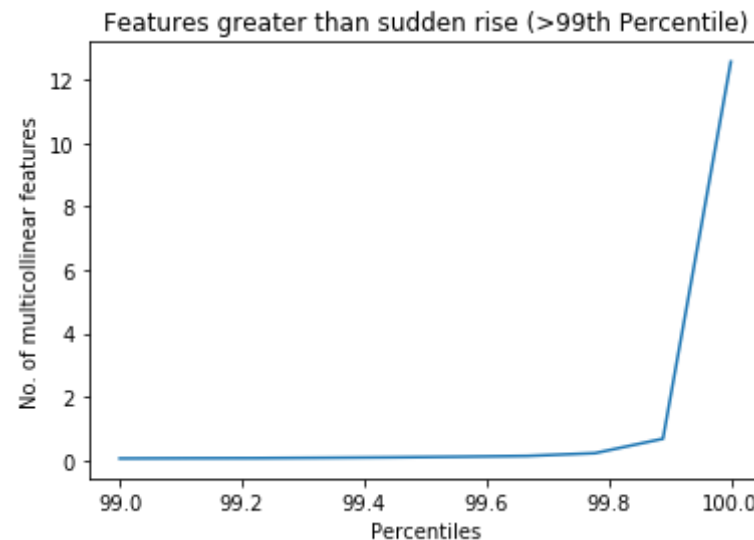
```
gt99= np.linspace(99,100,10)
tab = PrettyTable()
tab.field_names = ['Percentile ', 'Percentile Value']

for i in gt99:
    tab.add_row([i, np.percentile(W_percent_change_vector,i)])

print(tab)
plt.plot(gt99,np.percentile(W_percent_change_vector,gt99) );
plt.xlabel('Percentiles')
plt.ylabel('No. of multicollinear features')
plt.title("Features greater than sudden rise (>99th Percentile)")
del (tab)
```

Percentile	Percentile Value
99	0
99.1	0
99.2	0
99.3	0
99.4	0
99.5	0
99.6	0
99.7	0
99.8	0
99.9	0
100	12.5

99.0	0.047399636184165865
99.11111111111111	0.052915605599829704
99.22222222222223	0.05635737937635935
99.33333333333333	0.06811388084312188
99.44444444444444	0.08260693710378353
99.55555555555556	0.10265356775837843
99.66666666666667	0.12710099736555297
99.77777777777777	0.21553890490009822
99.88888888888889	0.6702054059627296
100.0	12.572261483067537



Observation: In this case the threshold value is the percentile **99.88**

```
In [90]: # We will list down the features names which are greater than this index value of threshold.
# We will print top 10 features by sorting values in descending values of the % changed vector
```



```

feature_names = bow_vectorizer.get_feature_names() # get feature names
weight_vector_values = model_bow_l2_noisy.coef_      # get weight values

# sorting change in percent weight vector in descending order i.e. features that have undergone maximum collinearity
# and selecting top 10 features out of it
mulc_features = np.argsort(W_percent_change_vector)[:,-1][0,0:10]

# Top 10 features are
print("Printing 10 features:")

#tab = PrettyTable()
#tab.field_names = ['Feature name', 'Weight vector value']

for f in mulc_features:
    print(feature_names[f])
    #tab.add_row([feature_names[f], weight_vector_values[0,f]])

#print(tab)
#del(tab)

```

```

Printing 10 features:
minimal
free dairy
include
bowl
started using
quantity
good chocolate
not amazon
drink good
organic ingredients

```

Observation

1. We performed the perturbation test to find multicollinearity.
2. Here, we can see that very few features are collinear. Less than 0.02% features are multicollinear

3. We can conclude that this model is not affected by multicollinearity of features

feature importance

[5.1.3] Feature Importance on BOW, SET 1

```
In [91]: model_bow_l2.classes_
```

```
Out[91]: array([0, 1], dtype=int64)
```

[5.1.3.1] Top 10 important features of positive class from SET 1

[5.1.3.2] Top 10 important features of negative class from SET 1

```
In [92]: most_informative_feature_for_binary_classification(bow_vectorizer, model_bow_l2)
```

Class	Coefficient (Importance)	Feature Name
0	0.18118765370558954	disappointed
0	0.1637476142766516	worst
0	0.1450317900116653	not recommend
0	0.1437890099604434	not good
0	0.14138196686305474	not worth
0	0.1394686428948509	not
0	0.13404722441563235	terrible
0	0.13008740467603347	awful
0	0.12911041632708453	not buy
0	0.12699445630766656	disappointing

Class	Coefficient (Importance)	Feature Name
-------	--------------------------	--------------

1	0.3205556750029913	great
1	0.2546191552577402	good
1	0.24534971037457565	delicious
1	0.22623311139831281	best
1	0.22139260541932348	love
1	0.18358386635904556	loves
1	0.17424257182496358	perfect
1	0.16690705041432588	excellent
1	0.14563081152530527	wonderful
1	0.14435945555081128	tasty

[4.2] Bi-Grams and n-Grams.

```
In [0]: #bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature\_extraction.text.CountVectorizer.html

# you can choose these numebrs min_df=10, max_features=5000, of your choice
#count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
#final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
#print("the type of count vectorizer ",type(final_bigram_counts))
#print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
#print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

[4.3] TF-IDF

```
In [93]: # ss
from sklearn.feature_extraction.text import TfidfVectorizer
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_tfidf = tf_idf_vect.transform(X_train)
#X_cv_tfidf = tf_idf_vect.transform(X_cv)
X_test_tfidf = tf_idf_vect.transform(X_test)

print("After vectorizations")
print(X_train_tfidf.shape, y_train.shape)
#print(X_cv_tfidf.shape, y_cv.shape)
print(X_test_tfidf.shape, y_test.shape)
print("="*100)
```

```
After vectorizations
(61441, 36173) (61441,)
(26332, 36173) (26332,)
=====
=====
```


```
In [94]: print("the type of count vectorizer ",type(X_train_tfidf))
print("the shape of cut text TFIDF vectorizer ",X_train_tfidf.get_shape
())
print("the number of unique words: ", X_train_tfidf.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of cut text TFIDF vectorizer (61441, 36173)
the number of unique words: 36173
```

[5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
In [95]: # Get the best hyperparameter using GridSearchCV with penalty l1 and cv
= 5
```

```
best_estimator_tfidf_l1 = get_best_hyperparameter_C(tf_idf_vect, X_train_tfidf, X_test_tfidf, y_train, y_test, penalty_l = 'l1' )
```




Mean CV Score	Std CV Score	Param
0.5	0.0	{'C': 0.0001}
0.5	0.0	{'C': 0.001}
0.664	0.01405	{'C': 0.01}
0.886	0.01464	{'C': 0.1}
0.952	0.00962	{'C': 1}
0.949	0.01358	{'C': 10}
0.942	0.01555	{'C': 100}
0.94	0.0167	{'C': 1000}
0.932	0.01891	{'C': 10000}

```
The best estimator:LogisticRegression(C=1, class_weight=None, dual=False,
fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l1', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

The best score is:0.9516803305794055
The best value of C is:{'C': 1}
Mean Score: 0.9604442465849716

```
In [96]: # Fitting the model with the best hyperparameter
model_tfidf_l1 = LogisticRegression(C= list(best_estimator_tfidf_l1.val
ues())[0] ,penalty = 'l1')
model_tfidf_l1.fit(X_train_tfidf,y_train)
y_pred = model_tfidf_l1.predict(X_test_tfidf)
```

```
In [97]: # AUC- ROC plot
auc_train_tfidf_l1, auc_test_tfidf_l1 = plot_auc(model_tfidf_l1, X_train_tfidf, X_test_tfidf)
```



train AUC: 0.9672404879416822
test AUC: 0.9604436620765118

```
In [98]: # Confusion Matrix
print_confusion_matrix(model_tfidf_l1, X_train_tfidf, X_test_tfidf)

*****Train confusion matrix*****
```

```
[[ 6578  3046]
 [  871 50946]]
```

```
*****Test confusion matrix*****
```


```
[[ 3056  1501]
 [  496 21279]]
```

```
In [99]: # Heatmap Confusion Matrix
plt.figure(1)
plt.figure(figsize=(15, 4))

plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_tfidf_l1.predict(X_test_tfidf))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Test Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_tfidf_l1.predict(X_train_tfidf))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Train Set Confusion Matrix');

<Figure size 432x288 with 0 Axes>
```




Observation

1. For the TFIDF vectorizer, we calculated $C = 1$ using GridSearchCV with $cv = 5$ and with penalty l1.
2. We got train AUC: 0.9672404879416822 and test AUC: 0.9604436620765118
3. Using the confusion matrix, we can say that our model correctly predicted 21279 positive reviews and 3055 negative reviews.
4. The model incorrectly classified 496 negative reviews and 1501 positive reviews.
5. The True Postive Rate is 97.84 and the True Negative Rate is 67.03
6. The accuracy of the model is 92.50

[5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

```
In [100]: # Get the best hyperparameter
best_estimator_tfidf_l2 = get_best_hyperparameter_C(tf_idf_vect, X_train_tfidf, X_test_tfidf, y_train, y_test, penalty_l = 'l2' )
```


Mean CV Score	Std CV Score	Param
0.712	0.02389	{'C': 0.0001}
0.913	0.01532	{'C': 0.001}
0.912	0.01319	{'C': 0.01}
0.935	0.01353	{'C': 0.1}
0.957	0.01182	{'C': 1}
0.958	0.01198	{'C': 10}
0.951	0.01371	{'C': 100}
0.947	0.01465	{'C': 1000}
0.945	0.01516	{'C': 10000}

The best estimator: LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)

The best score is: 0.95818493762303
The best value of C is: {'C': 10}
Mean Score: 0.964725771053579

```
In [101]: # Fit the model with the best hyperparameter
model_tfidf_l2 = LogisticRegression(C= list(best_estimator_tfidf_l2.val
ues())[0], penalty = 'l2')
model_tfidf_l2.fit(X_train_tfidf,y_train)
y_pred = model_tfidf_l2.predict(X_test_tfidf)
```

```
In [102]: # AUC- ROC plot
auc_train_tfidf_l2, auc_test_tfidf_l2 = plot_auc(model_tfidf_l2, X_train_tfidf, X_test_tfidf)
```




train AUC: 0.998516672221631
test AUC: 0.964725771053579

```
In [103]: # Confusion matrix
print_confusion_matrix(model_tfidf_l2, X_train_tfidf, X_test_tfidf)

*****Train confusion matrix*****
[[ 8946   678]
 [  173 51644]]
```

```
*****Test confusion matrix*****  
[[ 3248  1309]  
 [  566 21209]]
```

```
In [104]: # Heatmap confusion Matrix  
plt.figure(1)  
plt.figure(figsize=(15, 4))  
  
plt.subplot(121) # Test confusion matrix  
cnf_matrix = confusion_matrix(y_test, model_tfidf_l2.predict(X_test_tfidf))  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Test Set Confusion Matrix');  
  
plt.subplot(122) # Train Confusion matrix  
cnf_matrix = confusion_matrix(y_train, model_tfidf_l2.predict(X_train_tfidf))  
np.set_printoptions(precision=2)  
class_names = ['Negative', 'Positive']  
# Plot non-normalized confusion matrix  
#plt.figure()  
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Train Set Confusion Matrix');  
  
<Figure size 432x288 with 0 Axes>
```



Observation

1. For the TF-IDF vectorizer, we calculated $C = 10$ using GridSearchCV with $cv = 5$ and with penalty l2.
2. We got train AUC: 0.9985166762221631 and test AUC: 0.964725771053579
3. Using the confusion matrix, we can say that our model correctly predicted 21209 positive reviews and 3248 negative reviews.
4. The model incorrectly classified 566 negative reviews and 1309 positive reviews.
5. The True Positive Rate is 98.47 and the True Negative Rate is 60.21
6. The accuracy of the model is 91.85

[5.2.3] Feature Importance on TFIDF, SET 2

[5.2.3.1] Top 10 important features of positive class from SET 2

[5.2.3.2] Top 10 important features of negative class from SET 2

```
In [105]: most_informative_feature_for_binary_classification(tf_idf_vect, model_tfidf_l2)
```

```
+-----+-----+-----+-----+
```

Class	Coefficient (Importance)	Feature Name
0	14.05633689502798	worst
0	11.74311885235004	not worth
0	11.712297104959422	disappointed
0	11.553933190987355	terrible
0	11.31806567269511	not good
0	11.294343181493128	not recommend
0	11.235079122127873	two stars
0	10.916559085516361	awful
0	10.615169158405667	disappointing
0	9.237292164583296	disappointment

Class	Coefficient (Importance)	Feature Name
1	16.229861601293226	great
1	14.71476530913575	delicious
1	12.67482046352735	not disappointed
1	12.54874612418282	good
1	12.181395785494692	best
1	11.90106615797133	perfect
1	11.111630609138262	loves
1	10.525449587060917	excellent
1	10.043258129638604	amazing
1	9.770806515013575	wonderful

[4.4] Word2Vec

```
In [106]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence_train=[]
for sentence in X_train:
    list_of_sentence_train.append(sentence.split())
```

```
In [108]: print(list_of_sentence_train[0])

['bought', 'apartment', 'infested', 'fruit', 'flies', 'hours', 'trap',
'attracted', 'many', 'flies', 'within', 'days', 'practically', 'gone',
'may', 'not', 'long', 'term', 'solution', 'flies', 'driving', 'crazy',
'consider', 'buying', 'one', 'caution', 'surface', 'sticky', 'try', 'av
oid', 'touching']
```

```
In [109]: is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence_train,min_count=5,size=50, work
ers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors
-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")

[('fantastic', 0.8422129154205322), ('good', 0.8153432011604309), ('excellent', 0.813062310218811), ('awesome', 0.8104643821716309), ('wonderful', 0.8010066747665405), ('terrific', 0.7817002534866333), ('perfect', 0.767562747001648), ('amazing', 0.7452276945114136), ('fabulous', 0.7190866470336914), ('decent', 0.7043147087097168)]
=====
[('greatest', 0.7702896595001221), ('best', 0.7596640586853027), ('tastiest', 0.7124402523040771), ('nastiest', 0.6849291920661926), ('coolest', 0.6719734072685242), ('closest', 0.6640547513961792), ('toughest',
```

```
0.6579124927520752), ('disgusting', 0.6277365684509277), ('smoothest',  
0.6193704605102539), ('ive', 0.6013782024383545)]
```

```
In [110]: w2v_words = list(w2v_model.wv.vocab)  
print("number of words that occurred minimum 5 times ", len(w2v_words))  
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 14799  
sample words ['bought', 'apartment', 'infested', 'fruit', 'flies', 'ho  
urs', 'trap', 'attracted', 'many', 'within', 'days', 'practically', 'go  
ne', 'may', 'not', 'long', 'term', 'solution', 'driving', 'crazy', 'con  
sider', 'buying', 'one', 'caution', 'surface', 'sticky', 'try', 'avoi  
d', 'touching', 'really', 'good', 'idea', 'final', 'product', 'outstand  
ing', 'use', 'car', 'window', 'everybody', 'asks', 'made', 'two', 'thum  
bs', 'received', 'shipment', 'could', 'hardly', 'wait', 'love', 'call']
```

Converting train text data

```
In [111]: # average Word2Vec  
# compute average word2vec for each review.  
sent_vectors_train = []; # the avg-w2v for each sentence/review is stor  
ed in this list  
for sent in tqdm(list_of_sentence_train): # for each review/sentence  
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, yo  
u might need to change this to 300 if you use google's w2v  
    cnt_words = 0; # num of words with a valid vector in the sentence/re  
view  
    for word in sent: # for each word in a review/sentence  
        if word in w2v_words:  
            vec = w2v_model.wv[word]  
            sent_vec += vec  
            cnt_words += 1  
    if cnt_words != 0:  
        sent_vec /= cnt_words  
    sent_vectors_train.append(sent_vec)  
sent_vectors_train = np.array(sent_vectors_train)  
print(sent_vectors_train.shape)  
print(sent_vectors_train[0])
```

```
(61441, 50)
[ 0.26 -0.27  0.11  0.07 -0.18  0.39 -0.26 -0.31  0.17 -0.01  0.11  0.4
5    0.32  0.34 -0.19 -0.17  0.37  0.16 -0.03  0.3   0.21  0.57 -0.26 -0.3
  0.73  0.11 -0.03  0.08  0.11  0.47  0.11 -0.14 -0.19  0.14 -0.36 -0.3
8    -0.04  0.02  0.57 -0.16  0.78  0.28 -0.42 -0.04  0.11 -0.41 -0.28 -0.2
1    -0.74 -0.19]
```

```
In [112]: i=0
list_of_sentence_test=[]
for sentence in X_test:
    list_of_sentence_test.append(sentence.split())
```



```

        fit_params=None, iid='warn', n_jobs=None,
        param_grid=[{'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 1
0000]}],
        pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
        scoring='roc_auc', verbose=0)

```

```
In [131]: best_estimator_w2v_l1 = gs_obj.best_params_
```

```
In [132]: list(best_estimator_w2v_l1.values())[0]
```

```
Out[132]: 100
```

```
In [140]: train_auc= gs_obj.cv_results_['mean_train_score']
train_auc_std= gs_obj.cv_results_['std_train_score']
cv_auc = gs_obj.cv_results_['mean_test_score']
cv_auc_std= gs_obj.cv_results_['std_test_score']

# draws the error plot

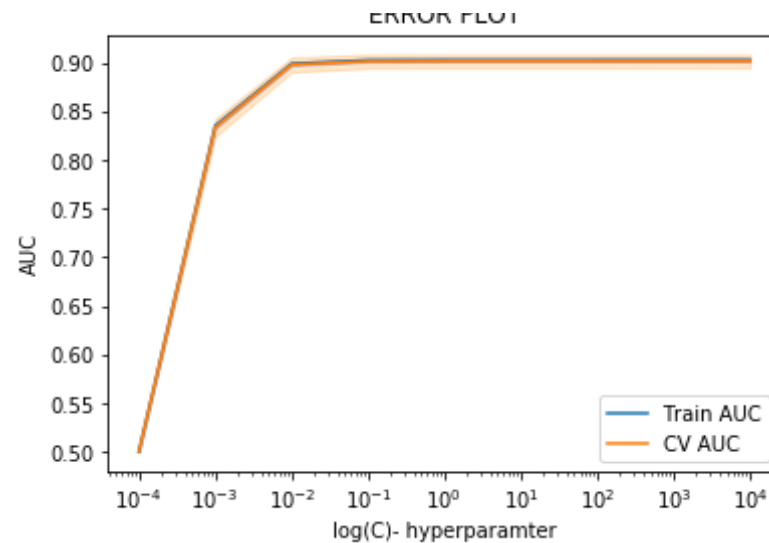
plt.plot(alpha, train_auc, label='Train AUC')

# this code is copied from here: https://stackoverflow.com/a/48803361/4
084039
plt.gca().fill_between(alpha,train_auc - train_auc_std,train_auc + train
_auc_std,alpha=0.2,color='darkblue')

plt.plot(alpha, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4
084039
plt.gca().fill_between(alpha,cv_auc - cv_auc_std,cv_auc + cv_auc_std,al
pha=0.2,color='darkorange')
plt.legend()
plt.xlabel("log(C) - hyperparamter")
plt.xscale('log')
plt.ylabel("AUC")
plt.title("ERROR PLOT")
plt.show()

```

ERROR PLOT



```
In [141]: # Results of the gs object
# Code https://stackoverflow.com/questions/42793254/what-replaces-grids
# earchcv-grid-scores-in-scikit#answer-42800056
means = gs_obj.cv_results_['mean_test_score']
stds = gs_obj.cv_results_['std_test_score']

t1 = PrettyTable()
t1.field_names = ['Mean CV Score', 'Std CV Score', 'Param']

for mean, std, params in zip(means, stds, gs_obj.cv_results_['params']):
    t1.add_row([round(mean, 3), round(std * 2, 5), params])

print(t1)

print("\nThe best estimator:{}".format(gs_obj.best_estimator_))
print("\nThe best score is:{}".format(gs_obj.best_score_))
print("The best value of C is:{}".format(gs_obj.best_params_))

# Returns the mean accuracy on the given test data and labels.
print("Mean Score: {}".format(gs_obj.score(sent_vectors_test, y_test)))
```


Mean CV Score	Std CV Score	Param
0.5	0.0	{'C': 0.0001}
0.833	0.01844	{'C': 0.001}
0.898	0.01507	{'C': 0.01}
0.901	0.0144	{'C': 0.1}
0.901	0.01406	{'C': 1}
0.901	0.01405	{'C': 10}
0.901	0.01404	{'C': 100}
0.901	0.01399	{'C': 1000}
0.901	0.01404	{'C': 10000}

The best estimator: LogisticRegression(C=1000, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l1', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)

The best score is: 0.901461054602771
The best value of C is: {'C': 1000}
Mean Score: 0.9015784802124991

```
In [142]: # Fitting the model with the best hyperparameter
model_avgw2v_l1 = LogisticRegression(C= list(best_estimator_w2v_l1.values())[0], penalty = 'l1')
model_avgw2v_l1.fit(sent_vectors_train, y_train)
y_pred = model_avgw2v_l1.predict(sent_vectors_test)
```

```
In [143]: # AUC - ROC plot
auc_train_avgw2v_l1, auc_test_avgw2v_l1 = plot_auc(model_avgw2v_l1, sent_vectors_train, sent_vectors_test)
```



train AUC: 0.902424731877006
test AUC: 0.9015779662481638

```
In [144]: # Confusion matrix
print_confusion_matrix(model_avgw2v_l1, sent_vectors_train, sent_vectors_test)
```

*****Train confusion matrix*****

```
[[ 4454  5170]
 [ 1726 50091]]
```

*****Test confusion matrix*****

```
[[ 2146  2411]
 [   788 20987]]
```

```
In [156]: # Heatmap confusion matrix
plt.figure(1)
plt.figure(figsize=(15, 4))
```

```
plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_avgw2v_l1.predict(sent_vect
```


```

ors_test))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
est Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_avgw2v_l1.predict(sent_vec
tors_train))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
rain Set Confusion Matrix');

```

<Figure size 432x288 with 0 Axes>



Observation

1. For the BoW vectorizer, we calculated $C = 1000$ using GridSearchCV with $cv = 5$ and with penalty l1.
2. We got train AUC: 0.902424731877006 and test AUC: 0.9015779662481638

3. Using the confusion matrix, we can say that our model correctly predicted 20987 positive reviews and 2146 negative reviews.
4. The model incorrectly classified 788 negative reviews and 2411 positive reviews.
5. The True Postive Rate is 96.41 and the True Negative Rate is 47.65
6. The accuracy of the model is 87.97

[5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

```
In [146]: # Get the best hyperparameter
tuned_parameters = [{ 'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1,
, 10**2, 10**3, 10**4]}]
alpha = [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10
**4] #k

# Using GridSearchCVSearchCV with 5 fold cv
gs_obj = GridSearchCV(LogisticRegression(penalty= 'l2'), tuned_paramete
rs, scoring = 'roc_auc', cv=5)

gs_obj.fit(sent_vectors_train, y_train)
```

```
Out[146]: GridSearchCV(cv=5, error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid=[{'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 1
0000]}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=0)
```

```
In [147]: train_auc= gs_obj.cv_results_['mean_train_score']
train_auc_std= gs_obj.cv_results_['std_train_score']
cv_auc = gs_obj.cv_results_['mean_test_score']
```

```


cv_auc_std= gs_obj.cv_results_['std_test_score']

# draws the error plot

plt.plot(alpha, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

plt.plot(alpha, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
plt.legend()
plt.xlabel("log(C): hyperparameter")
plt.xscale('log')
plt.ylabel("AUC")
plt.title("ERROR PLOT")
plt.show()

```




```
In [148]: # Results of the gs object
# Code https://stackoverflow.com/questions/42793254/what-replaces-grids
# earchcv-grid-scores-in-scikit#answer-42800056
means = gs_obj.cv_results_['mean_test_score']
stds = gs_obj.cv_results_['std_test_score']

t2 = PrettyTable()
t2.field_names = ['Mean CV Score', 'Std CV Score', 'Param']

for mean, std, params in zip(means, stds, gs_obj.cv_results_['params']):
    t2.add_row([round(mean, 3), round(std * 2, 5), params])

print(t2)

print("\nThe best estimator:{}".format(gs_obj.best_estimator_))
print("\nThe best score is:{}".format(gs_obj.best_score_))
print("The best value of C is:{}".format(gs_obj.best_params_))

# Returns the mean accuracy on the given test data and labels.
print("Mean Score: {}".format(gs_obj.score(sent_vectors_test, y_test)))
```

Mean CV Score	Std CV Score	Param
0.825	0.02433	{'C': 0.0001}
0.895	0.01698	{'C': 0.001}
0.901	0.01463	{'C': 0.01}
0.901	0.0141	{'C': 0.1}
0.901	0.01402	{'C': 1}
0.901	0.01401	{'C': 10}
0.901	0.01401	{'C': 100}
0.901	0.01402	{'C': 1000}
0.901	0.01402	{'C': 10000}

```
The best estimator:LogisticRegression(C=0.1, class_weight=None, dual=False,
fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
```

```
n_jobs=None, penalty='l2', random_state=None, solver='warn',  
tol=0.0001, verbose=0, warm_start=False)
```

The best score is:0.9014769955102173


The best value of C is:{'C': 0.1}

Mean Score: 0.9015764646660858

```
In [149]: best_estimator_w2v_l2 = gs_obj.best_params_
```

```
In [150]: # Fitting the Avg W2v vectorizer on LogisticRegression Model with pena  
lty = 'l2'  
model_avgw2v_l2 = LogisticRegression(C= list(best_estimator_w2v_l2.valu  
es())[0],penalty = 'l2')  
model_avgw2v_l2.fit(sent_vectors_train,y_train)  
y_pred = model_avgw2v_l2.predict(sent_vectors_test)
```

```
In [151]: # AUC- ROC plot  
auc_train_avgw2v_l2, auc_test_avgw2v_l2 = plot_auc(model_avgw2v_l2, sen  
t_vectors_train, sent_vectors_test)
```



train AUC: 0.9024168652161337

test AUC: 0.9015764646660858

```
In [152]: # Confusion Matrix
print_confusion_matrix(model_avgw2v_l2, sent_vectors_train, sent_vectors_test)

*****Train confusion matrix*****
[[ 4410  5214]
 [ 1688 50129]]


*****Test confusion matrix*****
[[ 2113  2444]
 [   768 21007]]
```

```
In [155]: # Heatmap confusion matrix
plt.figure(1)
plt.figure(figsize=(15, 4))

plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_avgw2v_l2.predict(sent_vectors_test))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
# plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Test Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_avgw2v_l2.predict(sent_vectors_train))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
# plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='Train Set Confusion Matrix');
```

<Figure size 432x288 with 0 Axes>



Observation

1. For the BoW vectorizer, we calculated $C = 0.1$ using GridSearchCV with $cv = 5$ and with penalty l2.
2. We got train AUC: 0.9024168652161337 and test AUC: 0.9015764646660858
3. Using the confusion matrix, we can say that our model correctly predicted 21007 positive reviews and 2113 negative reviews.
4. The model incorrectly classified 768 negative reviews and 2444 positive reviews.
5. The True Positive Rate is 96.42 and the True Negative Rate is 47.59
6. The accuracy of the model is 87.97

[4.4.1.2] TFIDF weighted W2v

[5.4] Logistic Regression on TFIDF W2V, SET 4

```
In [114]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
X_train_tf_idf_w2v = model.fit_transform(X_train)
```


```
X_test_tf_idf_w2v = model.transform(X_test)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [115]: # TF-IDF weighted Word2Vec for sentences in X_train
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and ce
ll_val = tfidf

tfidf_sent_vectors_train = []; # the tfidf-w2v for each sentence/review
is stored in this list
row=0;
for sent in tqdm(list_of_sentence_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/r
review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole courpus
# sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_train.append(sent_vec)
    row += 1
```

```
100%|███████████████████████████████████████████████████████████████████████████  
███████████████████████████████████████████████████████████████████████████████| 61441/61441  
[27:08<00:00, 37.73it/s]
```

```
In [116]: # TF-IDF weighted Word2Vec for sentences in X_test
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final tf idf is the sparse matrix with row= sentence, col=word and ce
```

Mean CV Score	Std CV Score	Param
0.5	0.0	{'C': 0.0001}
0.801	0.01918	{'C': 0.001}
0.872	0.01642	{'C': 0.01}
0.877	0.01663	{'C': 0.1}
0.877	0.0166	{'C': 1}
0.877	0.01661	{'C': 10}
0.877	0.01661	{'C': 100}
0.877	0.01661	{'C': 1000}
0.877	0.01661	{'C': 10000}


The best estimator: LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l1', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)

The best score is: 0.8769711699379582

The best value of C is:{'C': 1}
Mean Score: 0.8734700629631506

```
In [118]: # Fitting the TFIDF - weighted W2V vectorizer on LogisticRegression Model
model_tfidf2v_l1 = LogisticRegression(C= list(best_estimator_tfidf2v_l1.values())[0], penalty = 'l1')
model_tfidf2v_l1.fit(tfidf_sent_vectors_train,y_train)
y_pred = model_tfidf2v_l1.predict(tfidf_sent_vectors_test)
```

```
In [119]: # AUC- ROC plot
auc_train_tfidf2v_l1, auc_test_tfidf2v_l1 = plot_auc(model_tfidf2v_l1, tfidf_sent_vectors_train, tfidf_sent_vectors_test)
```



train AUC: 0.8785712594988075
test AUC: 0.8734706575493425

```
In [120]: # Confusion Matrix
print_confusion_matrix(model_tfidf2v_l1, tfidf_sent_vectors_train, tfidf_sent_vectors_test)
```



```
*****Train confusion matrix*****
```

```
[[ 3636  5988]
 [ 1577 50240]]
```

```
*****Test confusion matrix*****
```


```
[[ 1731  2826]
 [   704 21071]]
```

```
In [121]: # Heatmap Confusion Matrix
plt.figure(1)
plt.figure(figsize=(15, 4))

plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_tfidfv2v_l1.predict(tfidf_s
ent_vectors_test))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
est Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_tfidfv2v_l1.predict(tfidf_
sent_vectors_train))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
rain Set Confusion Matrix');
```

```
<Figure size 432x288 with 0 Axes>
```




Observation

1. For the BoW vectorizer, we calculated $C = 1$ using GridSearchCV with $cv = 5$ and with penalty l1.
2. We got train AUC: 0.8785712594988075 and test AUC: 0.8734706575493425
3. Using the confusion matrix, we can say that our model correctly predicted 21071 positive reviews and 1731 negative reviews.
4. The model incorrectly classified 704 negative reviews and 2826 positive reviews.
5. The True Positive Rate is 96.60 and the True Negative Rate is 39.52
6. The accuracy of the model is 86.72

[5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

```
In [122]: best_estimator_tfidfw2v_l2 = get_best_hyperparameter_C(model, tfidf_sent_vectors_train, tfidf_sent_vectors_test, y_train, y_test, penalty='l2')
```



Mean CV Score	Std CV Score	Param
0.795	0.02675	{'C': 0.0001}
0.869	0.01828	{'C': 0.001}
0.876	0.01677	{'C': 0.01}
0.877	0.01661	{'C': 0.1}
0.877	0.01662	{'C': 1}
0.877	0.01661	{'C': 10}
0.877	0.01661	{'C': 100}
0.877	0.01661	{'C': 1000}
0.877	0.01661	{'C': 10000}

The best estimator: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)

The best score is: 0.8769926148015947


The best value of C is: {'C': 0.1}

Mean Score: 0.8735281610885159

train_auc = 0.8785701044652459

```
In [123]: # Fitting the TFIDF - weighted W2V vectorizer on LogisticRegression Model with C = and penalty = 'l2'
model_tfidf2v_l2 = LogisticRegression(C= list(best_estimator_tfidf2v_l2.values())[0],penalty = 'l2')
model_tfidf2v_l2.fit(tfidf_sent_vectors_train,y_train)
y_pred = model_tfidf2v_l2.predict(tfidf_sent_vectors_test)
```

```
In [124]: # AUC- ROC plot
auc_train_tfidf2v_l2, auc_test_tfidf2v_l2 = plot_auc(model_tfidf2v_l2, tfidf_sent_vectors_train, tfidf_sent_vectors_test)
```



train AUC: 0.8785701044652459
test AUC: 0.8735281610885159

```
In [125]: # Confusion Matrix
print_confusion_matrix(model_tfidf2v_l2, tfidf_sent_vectors_train, tfidf_sent_vectors_test)

*****Train confusion matrix*****
[[ 3607  6017]
```

```
[ 1548 50269]]


*****Test confusion matrix*****
[[ 1708  2849]
 [   690 21085]]
```

```
In [126]: # Heatmap Confusion Matrix
plt.figure(1)
plt.figure(figsize=(15, 4))

plt.subplot(121) # Test confusion matrix
cnf_matrix = confusion_matrix(y_test, model_tfidfv2v_l2.predict(tfidf_s
ent_vectors_test))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
est Set Confusion Matrix');

plt.subplot(122) # Train Confusion matrix
cnf_matrix = confusion_matrix(y_train, model_tfidfv2v_l2.predict(tfidf_
sent_vectors_train))
np.set_printoptions(precision=2)
class_names = ['Negative', 'Positive']
# Plot non-normalized confusion matrix
#plt.figure()
plot_confusion_matrix_heatmap(cnf_matrix, classes=class_names, title='T
rain Set Confusion Matrix');

<Figure size 432x288 with 0 Axes>
```



Observation

1. For the TFIDF-W2V vectorizer, we calculated $C = 0.1$ using GridSearchCV with $cv = 5$ and with penalty l2.
2. We got train AUC: 0.8785701044652459 and test AUC: 0.8735281610885159
3. Using the confusion matrix, we can say that our model correctly predicted 21085 positive reviews and 1702 negative reviews.
4. The model incorrectly classified 690 negative reviews and 2849 positive reviews.
5. The True Postive Rate is 96.6 and the True Negative Rate is 39.54
6. The accuracy of the model is 86.73

Conclusions

```
In [157]: C = PrettyTable()

C.field_names = ['Sr. No', 'Vectorizer', 'Norm', 'C', 'Train AUC', 'Test AUC']
C.add_row([1, 'BoW', 'L1', list(best_estimator_bow_l1.values())[0], auc_train_bow_l1, auc_test_bow_l1])
C.add_row([1, 'BoW', 'L2', list(best_estimator_bow_l2.values())[0], auc_train_bow_l2, auc_test_bow_l2])
```

```

C.add_row([1, 'TF_IDF', 'L1', list(best_estimator_tfidf_l1.values())[0], auc_train_tfidf_l1, auc_test_tfidf_l1])
C.add_row([1, 'TF_IDF', 'L2', list(best_estimator_tfidf_l2.values())[0], auc_train_tfidf_l2, auc_test_tfidf_l2])
C.add_row([1, 'Avg-W2V', 'L1', list(best_estimator_w2v_l1.values())[0], auc_train_avgw2v_l1, auc_test_avgw2v_l1])
C.add_row([1, 'Avg-W2V', 'L2', list(best_estimator_w2v_l2.values())[0], auc_train_avgw2v_l2, auc_test_avgw2v_l2])
C.add_row([1, 'TFIDF-W2V', 'L1', list(best_estimator_tfidfw2v_l1.values())[0], auc_train_tfidfw2v_l1, auc_test_tfidfw2v_l1])
C.add_row([1, 'TFIDF-W2V', 'L2', list(best_estimator_tfidfw2v_l2.values())[0], auc_train_tfidfw2v_l2, auc_test_tfidfw2v_l2])

print(C)
del C

```

```

+-----+-----+-----+-----+-----+-----+
-----+
| Sr. No | Vectorizer | Norm | C | Train AUC | Test AUC |
| UC |
+-----+-----+-----+-----+-----+-----+
-----+
| 1 | BoW | L1 | 0.01 | 0.9695167883807346 | 0.9695167883807346 |
| 1 | BoW | L2 | 0.001 | 0.99309054211035 | 0.9415796492294188 |
| 1 | TF_IDF | L1 | 1 | 0.9672404879416822 | 0.9604436620765118 |
| 1 | TF_IDF | L2 | 10 | 0.9985166762221631 | 0.9985166762221631 |
| 1 | Avg-W2V | L1 | 100 | 0.902424731877006 | 0.9015779662481638 |
| 1 | Avg-W2V | L2 | 0.1 | 0.9024168652161337 | 0.9015764646660858 |
| 1 | TFIDF-W2V | L1 | 1 | 0.8785712594988075 | 0.8734706575493425 |
| 1 | TFIDF-W2V | L2 | 0.1 | 0.8785701044652459 | 0.8735281610885159 |

```

Summary

1. We performed Logistic Regression with L1 and L2 normalization on BoW, TFIDF, Avg-W2V, TFIDF-WW2V on the Amazon Fine Food Reviews.
2. Made use of GridSearchCV to find the best value of C, the hyperparameter in logistic regression.
3. Performed Feature Engineering on the BoW model and found out the model slightly performed better.
4. Different vectors take on different C values. We saw C values being taken from 10^{-4} to 10^4
5. We obtained the colinear features using perturbation and found out that we have less than 0.02% of colinear features.
6. If dataset is balanced, we can get better TNR.